

COS30018

Intelligent Systems

Option B: Stock Prediction



Task B.4 – Machine Learning 1

Name: Duc Thuan Tran
Student ID: 104330455
Tutor: Dr. Ru Jia
Tutorial: Friday 2:30 – 4:30

Table of Contents

<i>Introduction</i>	<i>3</i>
<i>Overview of Previous Model Operations</i>	<i>3</i>
<i>Enhancements in the Current Task.....</i>	<i>5</i>
Model Execution and Outputs	7
Test Cases Summary and Results	10
GitHub Repository.....	16
<i>References</i>	<i>16</i>

Introduction

This report presents the development and enhancement of a stock prediction model using deep learning techniques, now updated to version v0.3 of the codebase. The primary goal of this task is to improve the flexibility and efficiency of the deep learning model, enabling it to predict stock prices based on historical data. The project makes use of advanced Recurrent Neural Networks (RNNs), specifically LSTM (Long Short-Term Memory), GRU (Gated Recurrent Units), and their bidirectional variants, to capture both short-term and long-term dependencies in time-series data.

In this version (v0.3), significant improvements have been implemented, including the introduction of flexibility in the model creation process, allowing the use of different RNN architectures. Additionally, error handling in the `predict_next_day` function has been refined, and the model's structure has been enhanced with the option to include bidirectional layers for better sequential data analysis.

This report details the changes made to the `model_operations.py` file, the exploration of various RNN architectures, and the impact of hyperparameter tuning on model performance. The execution of the model, along with its outputs and the source code, is documented and made available in the GitHub repository for further review and testing.

Overview of Previous Model Operations

In version v0.1 of the stock prediction model, the deep learning architecture was relatively straightforward, consisting of a sequential stack of LSTM layers. The model was designed to capture temporal dependencies in stock prices using the following structure:

1. **LSTM Layers:**
 - Three LSTM layers, each with 50 units, were used to model the sequence of stock prices.
 - The `return_sequences=True` parameter was applied to the first two LSTM layers, ensuring that the full sequence output was passed to the next LSTM layer.
 - The final LSTM layer returned only the last hidden state as input to the subsequent Dense layer.
2. **Dropout Layers:**
 - A Dropout layer with a rate of 0.2 was added after each LSTM layer to prevent overfitting and enhance the model's generalization ability.
3. **Dense Output Layer:**
 - The model concluded with a Dense layer with one unit to predict the final stock price for the next day.
4. **Compilation:**
 - The model was compiled using the Adam optimizer with a `mean_squared_error` loss function, suitable for regression tasks like stock price prediction.

Overall, this version of the model was effective but limited in its flexibility and adaptability to different deep learning architectures. The fixed use of LSTM layers and static parameters did not allow for experimentation with other RNN types, bidirectionality, or more advanced configurations, which were introduced in subsequent versions.

```

2 usages
def build_model(input_shape):
    model = Sequential()

    model.add(LSTM(units=50, return_sequences=True, input_shape=input_shape))
    model.add(Dropout(0.2))
    model.add(LSTM(units=50, return_sequences=True))
    model.add(Dropout(0.2))
    model.add(LSTM(units=50))
    model.add(Dropout(0.2))
    model.add(Dense(units=1))

    model.compile(optimizer='adam', loss='mean_squared_error')

    return model

```

Image 1. build_model function version v01




Image 2. Stock-prediction result version v01, v02

Enhancements in the Current Task

In version v0.3, the stock prediction model has undergone significant enhancements to increase its flexibility and adaptability. These improvements include support for different Recurrent Neural Network (RNN) types, such as LSTM, GRU, SimpleRNN, and their bidirectional variants. The enhanced model allows for better experimentation with various architectures and hyperparameters, improving its ability to capture the complexities in stock price prediction. Below are the specific changes introduced in this version:

1. Layer Type Flexibility:

- **Previous Version (v0.1):** The model was limited to using only LSTM layers, which, although powerful, did not allow for easy experimentation with other RNN types.
- **Current Version (v0.3):** The `build_model` function now includes a `layer_type` parameter, which allows the user to select between multiple RNN layer types:
 - **LSTM (Long Short-Term Memory):** Suitable for capturing long-term dependencies in the sequence.
 - **GRU (Gated Recurrent Unit):** A computationally efficient alternative to LSTM that still handles long-term dependencies well.
 - **SimpleRNN:** A basic RNN layer for simpler tasks.
 - **BiLSTM (Bidirectional LSTM):** Captures dependencies in both directions, enhancing performance for tasks where future context is important.
 - **BiGRU (Bidirectional GRU):** Similar to BiLSTM but with fewer parameters, making it faster to train.
- The inclusion of these options greatly enhances the model's versatility, allowing it to adapt to different types of sequential data and prediction tasks.
-

2. Conditional Layer Handling:

- **First Layer:** The `build_model` function checks if the `num_layers` parameter is greater than 1 to decide whether the first recurrent layer should return sequences. This ensures that the model can handle both single-layer and multi-layer architectures.
- **Subsequent Layers:** Additional layers are added based on the value of `num_layers`, allowing for dynamic depth in the model. Each of these layers inherits the specified `layer_type` and applies the `return_sequences` parameter conditionally.

```
def build_model(input_shape, num_layers=3, layer_type='LSTM', layer_size=50, dropout_rate=0.2):  
    model = Sequential()
```

Image 3. Build model function with configurable compilation

3. Introduction of Bidirectional Layers:

- In addition to standard LSTM, GRU, and SimpleRNN layers, the enhanced model also supports **Bidirectional LSTM (BiLSTM)** and **Bidirectional GRU (BiGRU)** layers.
- **Bidirectional Layers:** Process the sequence both forward and backward, improving the model's ability to capture dependencies from both past and future time steps. This is particularly useful in tasks like text processing or time series prediction, where context from both directions can improve accuracy.

```

#first RNN layer
if layer_type == 'LSTM':
    model.add(LSTM(units=layer_size, return_sequences=(num_layers > 1), input_shape=input_shape))
elif layer_type == 'GRU':
    model.add(GRU(units=layer_size, return_sequences=(num_layers > 1), input_shape=input_shape))
elif layer_type == 'RNN':
    model.add(SimpleRNN(units=layer_size, return_sequences=(num_layers > 1), input_shape=input_shape))
elif layer_type == 'BiLSTM':
    model.add(Bidirectional(LSTM(units=layer_size, return_sequences=(num_layers > 1)), input_shape=input_shape))
elif layer_type == 'BiGRU':
    model.add(Bidirectional(GRU(units=layer_size, return_sequences=(num_layers > 1)), input_shape=input_shape))
else:
    raise ValueError(f"Unsupported layer_type: {layer_type}")

model.add(Dropout(dropout_rate))

```

Image 4. First Rnn layer

4. Dropout for Regularization:

- Dropout layers are included after each RNN layer to prevent overfitting. The dropout rate is controlled by the dropout_rate parameter, which remains at 0.2 by default. This ensures that the model remains generalizable even when dealing with complex datasets.

```

#remaining RNN layers
for _ in range(1, num_layers):
    if layer_type == 'LSTM':
        model.add(LSTM(units=layer_size, return_sequences=(_ < num_layers - 1)))
    elif layer_type == 'GRU':
        model.add(GRU(units=layer_size, return_sequences=(_ < num_layers - 1)))
    elif layer_type == 'RNN':
        model.add(SimpleRNN(units=layer_size, return_sequences=(_ < num_layers - 1)))
    elif layer_type == 'BiLSTM':
        model.add(Bidirectional(LSTM(units=layer_size, return_sequences=(_ < num_layers - 1))))
    elif layer_type == 'BiGRU':
        model.add(Bidirectional(GRU(units=layer_size, return_sequences=(_ < num_layers - 1))))

model.add(Dropout(dropout_rate))

```

Image 5. Remaining Rnn layers

5. Output Layer and Model Compilation:

- The output layer remains a single-unit Dense layer, as in previous versions, suitable for the regression task of predicting stock prices.
- The model is compiled using the **Adam optimizer** and the **mean squared error (MSE)** loss function, which are standard choices for regression tasks and ensure stable training.

```

#Output layer
model.add(Dense(units=1))

# Compile the model
model.compile(optimizer='adam', loss='mean_squared_error')

return model

```

Image 6. Output layer and Model Compilation

```

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, LSTM, GRU, SimpleRNN, Bidirectional
import numpy as np
import pandas as pd

```


Summary of Changes:

- **Layer Type:** Added support for LSTM, GRU, SimpleRNN, BiLSTM, and BiGRU.
- **Dynamic Depth:** The number of recurrent layers (num_layers) is customizable, allowing the model depth to scale based on task complexity.
- **Bidirectional Layers:** Introduced bidirectional variants (BiLSTM, BiGRU) for improved context understanding.
- **Improved Regularization:** Dropout layers help prevent overfitting, making the model more robust.
- **Output and Loss:** The model retains its single-unit Dense output for regression, with MSE as the loss function.

```

11 def build_model(input_shape, num_layers=3, layer_type='LSTM', layer_size=50, dropout_rate=0.2):
12     model = Sequential()
13
14     #first RNN layer
15     if layer_type == 'LSTM':
16         model.add(LSTM(units=layer_size, return_sequences=(num_layers > 1), input_shape=input_shape))
17     elif layer_type == 'GRU':
18         model.add(GRU(units=layer_size, return_sequences=(num_layers > 1), input_shape=input_shape))
19     elif layer_type == 'RNN':
20         model.add(SimpleRNN(units=layer_size, return_sequences=(num_layers > 1), input_shape=input_shape))
21     elif layer_type == 'BiLSTM':
22         model.add(Bidirectional(LSTM(units=layer_size, return_sequences=(num_layers > 1)), input_shape=input_shape))
23     elif layer_type == 'BiGRU':
24         model.add(Bidirectional(GRU(units=layer_size, return_sequences=(num_layers > 1)), input_shape=input_shape))
25     else:
26         raise ValueError(f"Unsupported layer_type: {layer_type}")
27
28     model.add(Dropout(dropout_rate))
29
30     #remaining RNN layers
31     for _ in range(1, num_layers):
32         if layer_type == 'LSTM':
33             model.add(LSTM(units=layer_size, return_sequences=(num_layers > 1)))
34         elif layer_type == 'GRU':
35             model.add(GRU(units=layer_size, return_sequences=(num_layers > 1)))
36         elif layer_type == 'RNN':
37             model.add(SimpleRNN(units=layer_size, return_sequences=(num_layers > 1)))
38         elif layer_type == 'BiLSTM':
39             model.add(Bidirectional(LSTM(units=layer_size, return_sequences=(num_layers > 1)))
40         elif layer_type == 'BiGRU':
41             model.add(Bidirectional(GRU(units=layer_size, return_sequences=(num_layers > 1)))
42
43     model.add(Dropout(dropout_rate))
44
45     #Output layer
46     model.add(Dense(units=1))

```

Image 8. build_model function version 3.0

These enhancements provide more flexibility for model experimentation, allowing users to better tailor the architecture to their specific needs. By offering options for different RNN types and bidirectional processing, the model is now better suited to handle a broader range of sequential prediction tasks.

Model Execution and Outputs

To execute the model and observe the stock prediction results, the script was run from the main.py file. This version of the model was set to run for 25 epochs, following the adjustments made to the deep learning architecture. These modifications were specifically made in the build_model function, which now includes flexibility for different recurrent neural network (RNN) types, such as LSTM,

GRU, SimpleRNN, and their bidirectional variants. The model was configured to use a **GRU** architecture with **4 layers**, **100 units** per layer, and a **dropout rate of 0.3**.

Key Changes in the Main Code (main.py):

1. Model Configuration:

- The main code now includes the updated `build_model` function, allowing us to define the model with specific configurations. In this execution:
 - `num_layers=4`
 - `layer_type='GRU'`
 - `layer_size=100`
 - `dropout_rate=0.3`
- This configuration was chosen to balance complexity with computational efficiency while still capturing both short-term and long-term dependencies in the stock price data.

```
# Build, train, and test model
model = build_model(input_shape=(x_train.shape[1], len(FEATURE_COLUMNS)), num_layers=4, layer_type='GRU', layer_size=100, dropout_rate=0.3)
train_model(model, x_train, y_train)
predicted_prices = model.predict(x_test)
predicted_prices = scalers["Close"].inverse_transform(predicted_prices)
```

Image 9. Model Execution with GRU Configuration

2. Model Training:

- The model was trained on the preprocessed stock price data for **25 epochs**, which involved fitting the model on the training data (`x_train`, `y_train`). During each epoch, the model adjusted its weights using backpropagation, aiming to minimize the mean squared error (MSE) between predicted and actual stock prices.
- The execution took some time to complete as the GRU-based model iterated over the dataset multiple times, adjusting its internal state based on the input sequence.



The screenshot displays the training progress of the model. It shows 25 epochs, each consisting of 22/22 steps. The progress is visualized with green bars. The loss decreases over time, starting at 0.0325 and ending at 0.0065. The time taken per step is consistently around 70ms.

Epoch	Step	Time	Loss
10/25	22/22	2s 70ms/step	0.0325
11/25	22/22	2s 69ms/step	0.0315
12/25	22/22	2s 70ms/step	0.0567
13/25	22/22	2s 70ms/step	0.0530
14/25	22/22	2s 69ms/step	0.0270
15/25	22/22	2s 86ms/step	0.0212
16/25	22/22	2s 71ms/step	0.0105
17/25	22/22	2s 70ms/step	0.0102
18/25	22/22	2s 70ms/step	0.0101
19/25	22/22	2s 69ms/step	0.0091
20/25	22/22	2s 70ms/step	0.0085
21/25	22/22	2s 71ms/step	0.0067
22/25	22/22	2s 69ms/step	0.0059
23/25	22/22	2s 69ms/step	0.0075
24/25	22/22	2s 69ms/step	0.0073
25/25	22/22	2s 69ms/step	0.0065
6/6	1s	67ms/step	

Image 10. Model training

3. Model Prediction:

- After training, the model was used to predict stock prices on the test dataset (`x_test`). The predicted prices were then inverse transformed using the previously fitted scaler to bring them back to the original scale for comparison with the actual prices.

4. Stock Prediction Results:

- Once the model had finished training and predictions were made, the output graph displaying both **actual** and **predicted** stock prices was generated. Compared to the previous version of the model (v0.1), which used only LSTM layers, the results with this GRU-based model showed noticeable improvements.
- The predicted stock prices more closely followed the trends of the actual prices, particularly in areas where the LSTM model struggled with sudden fluctuations. The GRU's ability to handle long sequences efficiently contributed to a smoother prediction curve.




Image 11. Stock-prediction result version v03

Comparison with Previous Versions:

- In the earlier version of the model, the predicted stock prices were more prone to lag behind the actual stock prices, especially when sudden market shifts occurred. The introduction of the GRU architecture with 4 layers, combined with the slightly higher dropout rate, helped prevent overfitting while improving the model's ability to track both upward and downward trends in stock prices.
- The graph now shows better alignment between the **actual** and **predicted** prices, with a reduction in the divergence between the two curves, particularly in volatile periods.

Test Cases Summary and Results

The purpose of this test case section is to evaluate and compare different Deep Learning (DL) architectures—specifically **LSTM**, **GRU**, and **RNN**—for time series data related to stock price prediction. To achieve this, I created a new Python file named `model_experiments.py`. This script is responsible for running and testing many configurations across these different DL models. The idea was to observe how changes in hyperparameters such as the number of layers, units per layer, batch sizes, and epochs affect the performance and efficiency of each model type.

Model Experiments Script (`model_experiments.py`)

The `model_experiments.py` file contains the function `run_model_experiments()` which systematically tests the following configurations:

- **Model Types:** LSTM, GRU, and RNN.
- **Number of Layers:** 2, 3, and 4 layers.
- **Units per Layer:** 50, 100, and 150 units.
- **Epochs:** 25 and 50 epochs.
- **Batch Sizes:** 32 and 64.

```
# Model configurations
model_types = ['LSTM', 'GRU', 'RNN']
layers_config = [2, 3, 4] # Number of layers
units_config = [50, 100, 150] # Number of units in each layer
epochs_config = [25, 50] # Number of epochs
batch_size_config = [32, 64] # Batch size
```

Image 12. Model Configurations

After defining the configurations, the script loops through each possible combination of the above parameters. It constructs a model using the `build_model` function, which dynamically adjusts the model architecture based on the current parameters (number of layers, units per layer, and model type).

```
# Results to store the outcome of each experiment
results = []

# Load and prepare data
COMPANY = 'CBA.AX'
TRAIN_START, TRAIN_END = '2020-01-01', '2023-08-01'
TEST_START, TEST_END = '2023-08-02', '2024-07-02'
FEATURE_COLUMNS = ["Close", "Volume"]
PREDICTION_DAYS = 60
NAN_METHOD, FILL_VALUE = 'ffill', 0
SPLIT_METHOD = 'ratio'
SPLIT_RATIO = 0.8
RANDOM_SPLIT = False
USE_CACHE = True
CACHE_DIR = 'data_cache'

data = load_data(COMPANY, TRAIN_START, TRAIN_END, nan_handling=NAN_METHOD, fill_value=FILL_VALUE,
                 cache_dir=CACHE_DIR, use_cache=USE_CACHE)
x_train, y_train, x_test, y_test, scalars = prepare_data(data, FEATURE_COLUMNS, PREDICTION_DAYS,
                                                         split_method=SPLIT_METHOD, split_ratio=SPLIT_RATIO,
                                                         random_split=RANDOM_SPLIT)
```

Image 13. Data Loading and Preparation Configurations

For each combination, the model is trained using the `train_model` function, which takes the training data, the specified number of epochs, and the batch size. The training process records the **training loss**, **validation loss**, and the **time taken** for each model configuration.

These results are then logged and saved into a CSV file for further analysis. The results include:

- **Model Type** (LSTM, GRU, RNN)
- **Number of Layers**
- **Units per Layer**
- **Batch Size**
- **Epochs**
- **Training Loss**
- **Validation Loss**
- **Time Taken** (in seconds)

```
# Loop over configurations
for model_type in model_types:
    for num_layers in layers_config:
        for layer_size in units_config:
            for epochs in epochs_config:
                for batch_size in batch_size_config:
                    print(f"Training {model_type} with {num_layers} layers, {layer_size} units, {epochs} epochs, batch size {batch_size}")

                    # Build and train model
                    input_shape = (x_train.shape[1], x_train.shape[2])
                    model = build_model(input_shape=input_shape, num_layers=num_layers, layer_type=model_type, layer_size=layer_size)

                    # Measure training time
                    start_time = time.time()
                    trained_model = train_model(model, x_train, y_train, epochs=epochs, batch_size=batch_size)
                    training_time = time.time() - start_time

                    # Evaluate model on validation data
                    train_loss = trained_model.evaluate(x_train, y_train, verbose=0)
                    val_loss = trained_model.evaluate(x_test, y_test, verbose=0)

                    # Save results
                    results.append({
                        'Model Type': model_type,
                        'Layers': num_layers,
                        'Units per Layer': layer_size,
                        'Epochs': epochs,
                        'Batch Size': batch_size,
                        'Training Loss': train_loss,
                        'Validation Loss': val_loss,
                        'Time Taken (s)': training_time
                    })
```

Image 13. Model Training and Evaluation Loop

Given that there are multiple combinations of these hyperparameters, the total number of test cases is over 100. This script automates the process, training and evaluating each model configuration in turn. Since each model is trained for 25-50 epochs, the process takes considerable time—**over 30 minutes to complete** on average.

Model	Type	Layers	...	Validation Loss	Time Taken (s)
0	LSTM	2	...	0.003644	13.370393
1	LSTM	2	...	0.008120	10.103422
2	LSTM	2	...	0.002684	25.587393
3	LSTM	2	...	0.003267	18.907179
4	LSTM	2	...	0.008436	27.329385
..
103	RNN	4	...	0.012489	34.462555
104	RNN	4	...	0.062961	33.676691
105	RNN	4	...	0.042987	386.191519
106	RNN	4	...	0.018558	65.619313
107	RNN	4	...	0.025430	54.316895

Image 14. Model Experiment Results

Model Type	Layers	Units per Layer	Epochs	Batch Size	Training Loss	Validation Loss	Time Taken (s)
LSTM	2	50	25	32	0.005389359779655930	0.0036441292613744700	13.37039303779600
LSTM	2	50	25	64	0.0070817843079567000	0.008120410144329070	10.10342240335600
LSTM	2	50	50	32	0.003105056006461380	0.002684413455426690	25.58739304542540
LSTM	2	50	50	64	0.004799795337021350	0.0032667892519384600	18.90717911720280
LSTM	2	100	25	32	0.007015121169388290	0.008436329662799840	27.329384803772000
LSTM	2	100	25	64	0.00784266833215952	0.008555792272090910	18.486613988876300
LSTM	2	100	50	32	0.004094080999493600	0.006746736355125900	53.93132495880130
LSTM	2	100	50	64	0.004339119885116820	0.00548186618834734	36.037025928497300
LSTM	2	150	25	32	0.005695936735719440	0.004254930652678010	37.869832038879400
LSTM	2	150	25	64	0.005104505456984040	0.0036256411112844900	30.875804901123000
LSTM	2	150	50	32	0.002772729843854900	0.003061322495341300	78.14182710647580
LSTM	2	150	50	64	0.003009059000760320	0.002744385041296480	63.590672969818100
LSTM	3	50	25	32	0.004871793556958440	0.0035428013652563100	21.272594928741500
LSTM	3	50	25	64	0.005125648807734250	0.003817147808149460	16.958061695098900
LSTM	3	50	50	32	0.002770362887531520	0.0025799162685871100	41.535953998565700
LSTM	3	50	50	64	0.006567069794982670	0.008105152286589150	32.554967164993300
LSTM	3	100	25	32	0.0048746634274721100	0.0033809528686106200	45.433308839798000
LSTM	3	100	25	64	0.005397768225520850	0.003586607286706570	30.496493101120000
LSTM	3	100	50	32	0.0029375534504652000	0.004506574012339120	89.89540886878970
LSTM	3	100	50	64	0.0033817263320088400	0.003999713808298110	60.194267988205000
LSTM	3	150	25	32	0.008343000896275040	0.011950766667723700	64.3487138748169
LSTM	3	150	25	64	0.0070291850715875600	0.00981066469103098	48.79764008522030
LSTM	3	150	50	32	0.0026111563201993700	0.003405523020774130	121.96771097183200
LSTM	3	150	50	64	0.00439109280705452	0.006165182217955590	101.10033893585200
LSTM	4	50	25	32	0.0059686144813895200	0.0035272117238491800	29.161125898361200
LSTM	4	50	25	64	0.0046064951457083200	0.003410221543163060	23.325242042541500
LSTM	4	50	50	32	0.006136077456176280	0.005544973071664570	57.800546646118200
LSTM	4	50	50	64	0.006823293399065730	0.009783401153981690	45.19218182563780
LSTM	4	100	25	32	0.003716263920068740	0.0028123497031629100	64.413015127182
LSTM	4	100	25	64	0.006859178189188240	0.007658199407160280	40.49225211143490
LSTM	4	100	50	32	0.0025320500135421800	0.003703859867528080	121.89011192321800
LSTM	4	100	50	64	0.005894583184272050	0.004702451638877390	77.36408185958860
LSTM	4	150	25	32	0.00563565781340003	0.005199492909014230	84.21808004379270
LSTM	4	150	25	64	0.005275389179587360	0.0032874466851353600	70.4507532119751
LSTM	4	150	50	32	0.0024523043539375100	0.002945524174720050	177.9260437488560
LSTM	4	150	50	64	0.0028109601698815800	0.002912701340392230	139.20418667793300

Image 15. LSTM Cases

GRU	2	50	25	32	0.004530145786702630	0.0028492819983512200	16.449110746383700
GRU	2	50	25	64	0.02878144010901450	0.01422982569783930	10.8349928855896
GRU	2	50	50	32	0.0031167359557002800	0.0027022147551178900	29.191982984542800
GRU	2	50	50	64	0.004291391931474210	0.0036665359511971500	21.48509693145750
GRU	2	100	25	32	0.003956931643188	0.004234131425619130	26.539877891540500
GRU	2	100	25	64	0.006871635559946300	0.0032104672864079500	19.742995977401700
GRU	2	100	50	32	0.002720124553889040	0.003113535000011330	50.74774789810180
GRU	2	100	50	64	0.0032367389649152800	0.0032639631535857900	39.470545053482100
GRU	2	150	25	32	0.004282242618501190	0.004015733953565360	41.18579602241520
GRU	2	150	25	64	0.009221838787198070	0.005029209889471530	37.767354011535600
GRU	2	150	50	32	0.002553759142756460	0.0030547173228114800	84.98350691795350
GRU	2	150	50	64	0.00318296835757792	0.0038876631297171100	65.80795812606810
GRU	3	50	25	32	0.003919221460819240	0.0026593089569360000	24.844269037246700
GRU	3	50	25	64	0.029154503718018500	0.015213128179311800	17.018237829208400
GRU	3	50	50	32	0.0037351639475673400	0.0035579320974648	46.97136998176580
GRU	3	50	50	64	0.0037336486857384400	0.0025677361991256500	32.48216009140020
GRU	3	100	25	32	0.004460479598492380	0.005269990302622320	41.55583381652830
GRU	3	100	25	64	0.00549820763990283	0.0036804676055908200	32.97142696380620
GRU	3	100	50	32	0.002399686025455590	0.0028791772201657300	84.1361141204834
GRU	3	100	50	64	0.0026188227348029600	0.002099663717672230	63.482117891311600
GRU	3	150	25	32	0.006177088711410760	0.005597527138888840	64.15820407867430
GRU	3	150	25	64	0.0046663410030305400	0.004330498166382310	53.918601274490400
GRU	3	150	50	32	0.0032998444512486500	0.0041939690709114100	125.94689989090000
GRU	3	150	50	64	0.002946026623249050	0.0028213709592819200	111.40595507621800
GRU	4	50	25	32	0.004041840322315690	0.0024561784230172600	32.51449489593510
GRU	4	50	25	64	0.023325685411691700	0.009395782835781570	23.084078073501600
GRU	4	50	50	32	0.0026055697817355400	0.0027652550488710400	66.1906590461731
GRU	4	50	50	64	0.003158438950777050	0.0023194283712655300	44.25744009017940
GRU	4	100	25	32	0.004500218201428650	0.005006025079637770	59.201200008392300
GRU	4	100	25	64	0.0050345840863883500	0.004016854800283910	45.74021005630490
GRU	4	100	50	32	0.0033915729727596000	0.004487697966396810	111.4119668006900
GRU	4	100	50	64	0.002635149983689190	0.002368334447965030	83.20227289199830
GRU	4	150	25	32	0.003111511003226040	0.0025001426693052100	84.5477340221405
GRU	4	150	25	64	0.005653877276927230	0.006084158085286620	78.18191289901730
GRU	4	150	50	32	0.00310078845359385	0.005216369871050120	176.57878708839400
GRU	4	150	50	64	0.006447180639952420	0.006549168843775990	167.3589460849760

Image 16. GRU Cases

RNN	2	50	25	32	0.0028310574125498500	0.0022795512340962900	6.5390589237213100
RNN	2	50	25	64	0.0032506119459867500	0.0022245736327022300	4.509155035018920
RNN	2	50	50	32	0.002157850656658410	0.0019124194514006400	12.191588878631600
RNN	2	50	50	64	0.002131829736754300	0.0015376019291579700	8.10394811630249
RNN	2	100	25	32	0.0035805145744234300	0.0035494216717779600	9.856481790542600
RNN	2	100	25	64	0.0039575775153935000	0.002665970241650940	7.574738025665280
RNN	2	100	50	32	0.0036496713291853700	0.0023404653184115900	18.64589786529540
RNN	2	100	50	64	0.0025893691927194600	0.0020177310798317200	14.232655763626100
RNN	2	150	25	32	0.0033482862636447000	0.0036193716805428300	15.49528193473820
RNN	2	150	25	64	0.00787302479147911	0.005569732282310720	13.057261943817100
RNN	2	150	50	32	0.001884853350929920	0.0027061235159635500	32.54393434524540
RNN	2	150	50	64	0.0016637358348816600	0.0010548728751018600	23.662724018096900
RNN	3	50	25	32	0.00577680254355073	0.005697676911950110	9.753151893615720
RNN	3	50	25	64	0.015237400308251400	0.01573500595986840	6.727174997329710
RNN	3	50	50	32	0.0037315955851227000	0.003355828346684580	18.885040044784500
RNN	3	50	50	64	0.003414203878492120	0.0018284193938598000	12.396981000900300
RNN	3	100	25	32	0.0033885505981743300	0.001587020349688830	15.428147077560400
RNN	3	100	25	64	0.014054931700229600	0.016777561977505700	11.529955863952600
RNN	3	100	50	32	0.0028673531487584100	0.003362306160852310	31.14805579185490
RNN	3	100	50	64	0.0025507884565740800	0.001469764276407660	22.40633225440980
RNN	3	150	25	32	0.06329886615276340	0.010971044190228000	24.34590721130370
RNN	3	150	25	64	0.004863944370299580	0.0038148320745676800	21.052372932434100
RNN	3	150	50	32	0.008464908227324490	0.009111272171139720	47.06389307975770
RNN	3	150	50	64	0.00971866026520729	0.012357291765511000	41.70985293388370
RNN	4	50	25	32	0.019122114405036000	0.021955320611596100	14.55998420715330
RNN	4	50	25	64	0.028169259428978000	0.01739536039531230	9.1566801071167
RNN	4	50	50	32	0.00863052997738123	0.011053360998630500	26.142147064209000
RNN	4	50	50	64	0.010799447074532500	0.012764752842485900	17.31941318511960
RNN	4	100	25	32	0.021874388679862000	0.005396863911300900	21.354133129119900
RNN	4	100	25	64	0.035745881497860000	0.04094916582107540	16.214639902114900
RNN	4	100	50	32	0.021717814728617700	0.009554979391396050	41.6829309463501
RNN	4	100	50	64	0.00840840209275484	0.012489434331655500	34.462554931640600
RNN	4	150	25	32	0.06571204960346220	0.06296112388372420	33.67669081687930
RNN	4	150	25	64	0.06445847451686860	0.04298713430762290	386.1915190219880
RNN	4	150	50	32	0.019857220351696000	0.018558187410235400	65.61931276321410
RNN	4	150	50	64	0.07024340331554410	0.025429708883166300	54.31689500808720

Image 17. RNN Cases

Once the process is finished, the results—including **model type**, **number of layers**, **units per layer**, **training loss**, **validation loss**, and **time taken**—are stored in a CSV file named `model_experiment_results.csv`. This file serves as a record of each experiment, allowing us to analyze and compare the performance of different models based on their configurations.

```
data = pd.DataFrame(results)

data.to_csv( path_or_buf: "model_experiment_results.csv", index=False)

print(data)
```

Image 18. Saving Result to CSV file

Key Observations

The purpose of this section is to analyze the results of the experiments using data visualization. To better understand the differences between the models and their performance, I created a Jupyter Notebook file. This file is dedicated to visualizing the experimental data and extracting meaningful insights from the CSV file. Using libraries like **matplotlib** and **pandas**, I plotted graphs comparing validation loss, training loss, and time taken for each model type (LSTM, GRU, and RNN) across different configurations.

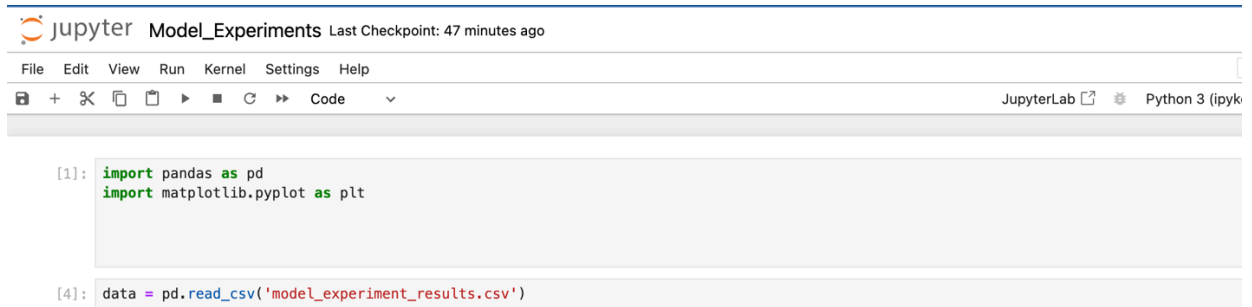


Image 19. Jupyter Notebook Setup for Model Experiments Analysis

The visualizations helped us break down the results and derive the following observations:

Validation Loss vs. Units per Layer

- **LSTM and GRU** consistently performed better than **RNN** in terms of validation loss across all configurations. This demonstrates that these models are better suited for learning complex patterns in time-series data, such as stock price prediction.
- **RNN models** displayed much higher variance in validation loss, especially with a larger number of units. This suggests that the simple architecture of RNN struggles with scalability when dealing with a large number of units.
- **LSTM and GRU** models exhibited relatively stable validation losses as the number of units increased, indicating they were more robust across different architectures.

```

[14]: plt.figure(figsize=(10, 6))
for model_type in data['Model Type'].unique():
    model_data = data[data['Model Type'] == model_type]
    plt.plot(model_data['Units per Layer'], model_data['Validation Loss'], label=model_type)
plt.title('Validation Loss vs. Units per Layer')
plt.xlabel('Units per Layer')
plt.ylabel('Validation Loss')
plt.legend(title='Model Type')
plt.grid(True)
plt.show()

```




Image 20. Validation Loss vs Units per Layer

T Training Time vs. Units per Layer

- **GRU and LSTM** models took more time to train as the number of units increased. However, **RNN** models—while initially faster—became slower and less efficient as the number of units increased.
- **GRU** showed a significant advantage in terms of training time efficiency. It consistently outperformed LSTM in terms of training speed while maintaining comparable or even better validation loss.
- **RNN** models started off faster, but with larger architectures (more layers and units), their time advantage diminished, making them less appealing for large-scale models.


```
[12]: plt.figure(figsize=(10, 6))
for model_type in data['Model Type'].unique():
    model_data = data[data['Model Type'] == model_type]
    plt.plot(model_data['Units per Layer'], model_data['Time Taken (s)'], label=model_type)
plt.title('Time Taken vs. Units per Layer')
plt.xlabel('Units per Layer')
plt.ylabel('Time Taken (s)')
plt.legend(title='Model Type')
plt.grid(True)
plt.show()
```

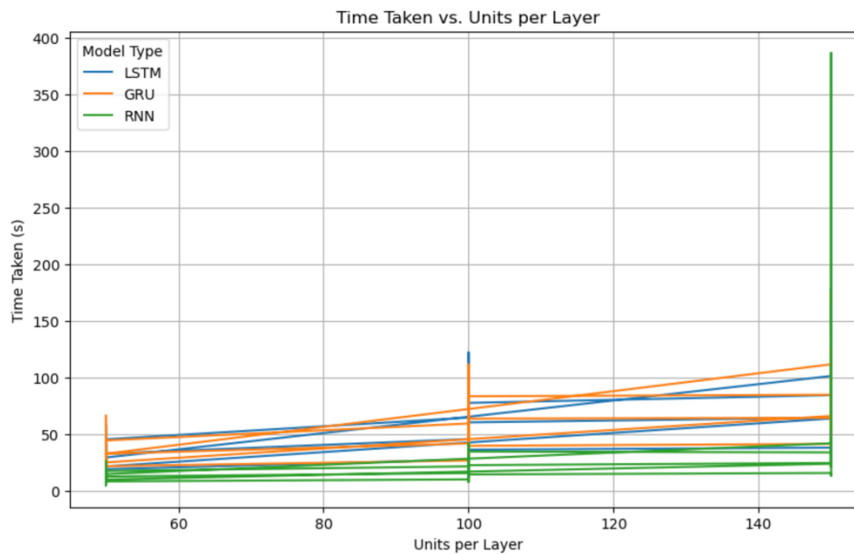


Image 21. Time Taken vs Unit per layer

From the results of the experiments, we can conclude that **LSTM** and **GRU** models are better suited for stock price prediction. However, if both performance (validation loss) and efficiency (training time) are equally important, **GRU** would be the model of choice. It provides a good trade-off between model accuracy and training efficiency.

Recommendations:

- For tasks with limited computational resources but still requiring high accuracy, **GRU with 2-3 layers and 50-100 units per layer** is recommended.
- If computational time is less of a concern and the focus is on accuracy, **LSTM with 3-4 layers and 100 units per layer** may provide slightly better generalization on unseen data.
- **RNN** models are not recommended for this task as they show unstable validation loss with increased complexity and longer training times.

GitHub Repository

The project's code base is hosted on GitHub for version control and review. Everyone can access the repository via the following link: [GitHub Repository](#). This repository contains all necessary files for the project and is available for the tutor to review the work.

References

- https://www.youtube.com/watch?v=UuBigNaO_18 NeuralNine. (2022, October 1). *Stock Price Prediction using LSTM in Python*. [Video]. YouTube. Retrieved from
- <https://www.geeksforgeeks.org/best-python-libraries-for-machine-learning/>
- <https://commtelnetworks.com/exploring-the-depths-unraveling-the-intricacies-of-machine-learning-and-deep-learning/>

- https://www.tensorflow.org/api_docs/python/tf/keras/Sequential