

1

Nature-inspired computing

- Nature has always served as a source of inspiration for engineers and scientists
- The best problem solver known in nature is:
 - **the (human) brain** that created “the wheel, New York, wars and so on” (after Douglas Adams’ Hitch-Hikers Guide)
 - **the evolution mechanism** that created the human brain (after Darwin’s Origin of Species)
- Answer 1 → neurocomputing
 - Week 6
- Answer 2 → evolutionary computing
 - Today + Week 8

2 / 39

2

Contents

- Motivations/applicable situations
- Basics of Evolutionary Computing (EC) Metaphor
- Basic scheme of an EA
- Basic Components:
 - Representation / Evaluation / Population / Parent Selection / Recombination / Mutation / Survivor Selection / Termination
- Examples : eight queens / knapsack
- Typical behaviours of EAs
- EC in context of global optimisation

3

Motivation

- Searching some search spaces with traditional search methods would be intractable. This is often true when states/candidate solutions have a large number of successors.
 - Example: Designing the surface of an aircraft.




Image source: <https://home.centaurosoft.com>

4

4

Evolutionary Computing (EC) – Main Idea

- **Adaptation** is Intelligence

(Nature)

→ Survival of the **Fittest** (aka. “natural selection”)

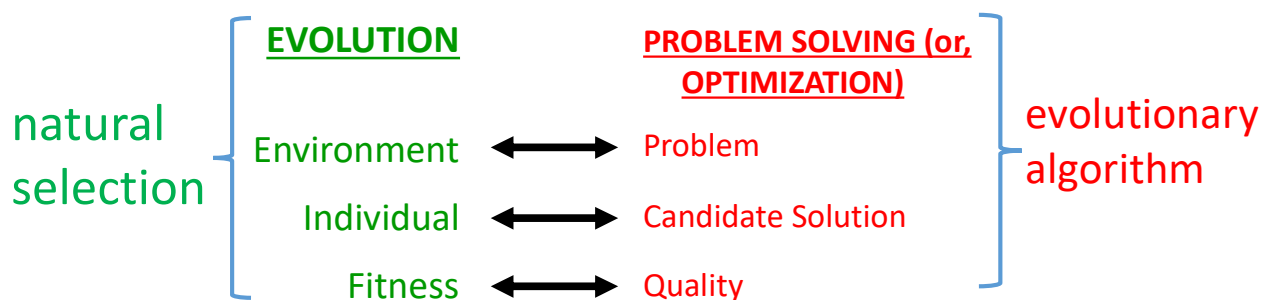
Darwin/Wallace’s theory: **Evolution** through *natural selection* of the *fittest* **individuals**

A process going through **multiple generations**

EC: How to use this idea for Optimization?

5

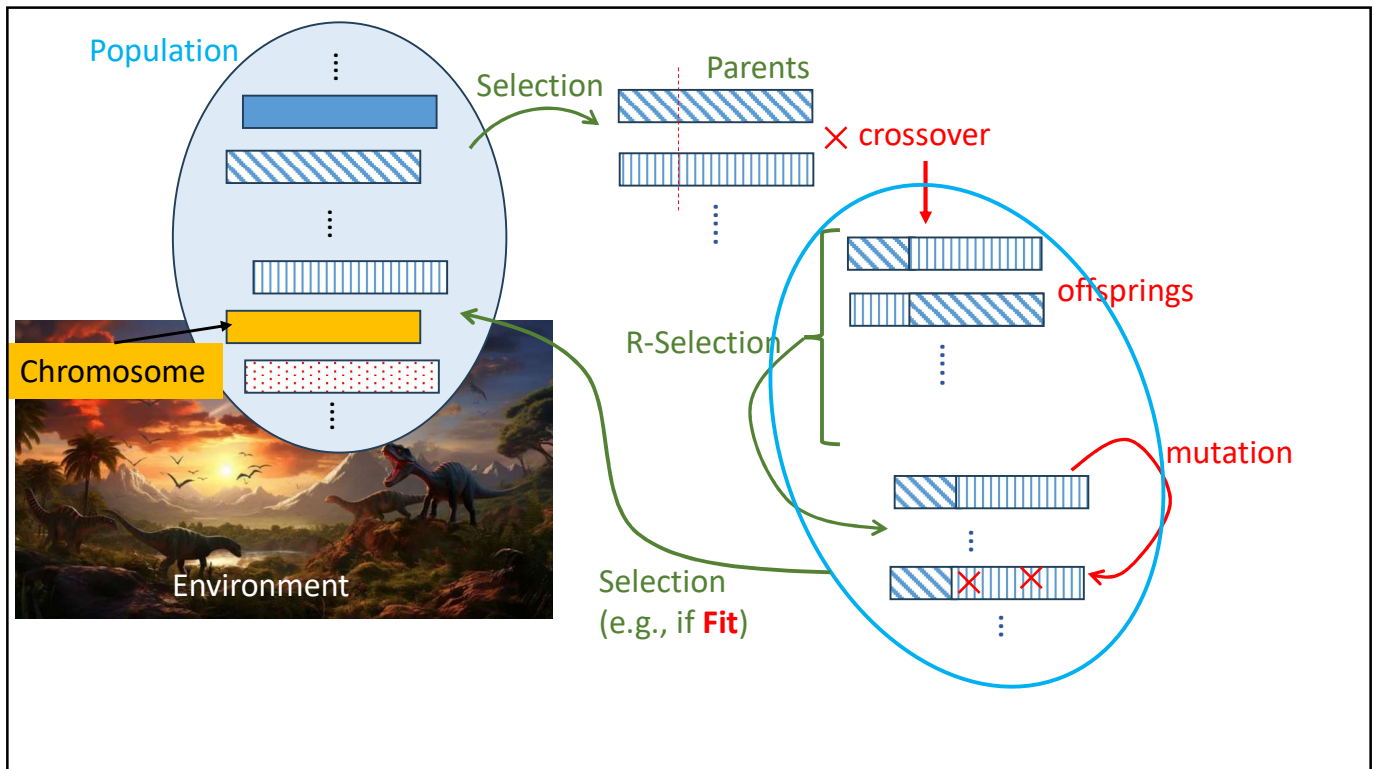
The Main EC Metaphor



Fitness → chances for survival and reproduction

Quality → chance for an existing solution to survive and seed new solutions

6



7

Basics of EC metaphor

- A **population of individuals** exists in an environment with limited resources
- **Competition** for those resources causes selection of those **fitter** individuals that are better adapted to the environment
- These individuals act as seeds for the generation of new individuals through recombination and mutation
- The new individuals have their fitness evaluated and compete (possibly also with parents) for survival.
- Over time **Natural selection** causes a rise in the fitness of the population

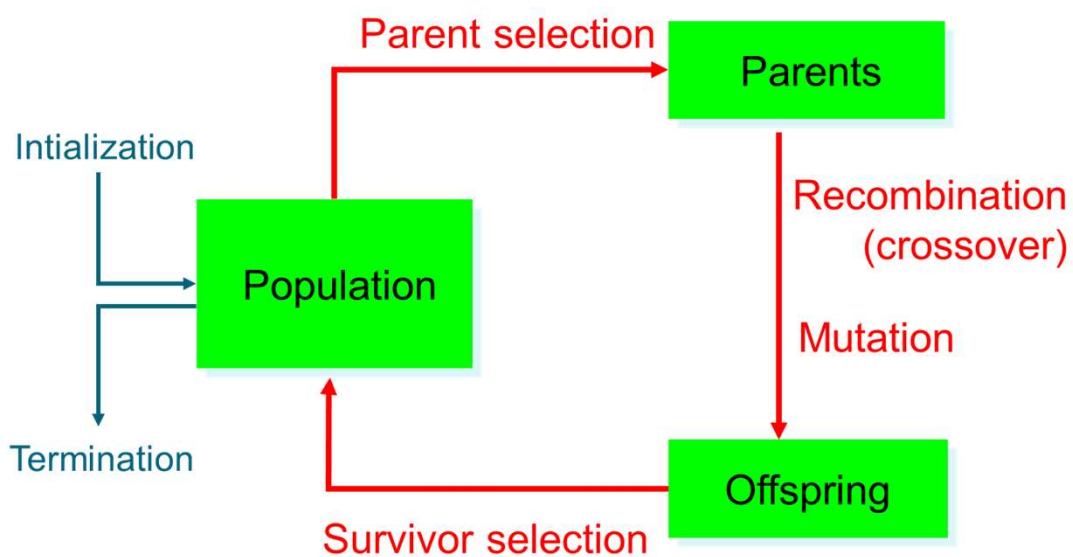
8

Basics of EC metaphor

- EAs fall into the category of “**generate and test**” algorithms
- They are stochastic, population-based algorithms
- Variation operators (recombination and mutation) create the necessary **diversity** and thereby facilitate novelty
- Selection reduces diversity and acts as a force pushing quality

9

General Scheme of EAs



10

Pseudo-code for typical EA

```

BEGIN
  INITIALISE population with random candidate solutions;
  EVALUATE each candidate;
  REPEAT UNTIL ( TERMINATION CONDITION is satisfied ) DO
    1 SELECT parents;
    2 RECOMBINE pairs of parents;
    3 MUTATE the resulting offspring;
    4 EVALUATE new candidates;
    5 SELECT individuals for the next generation;
  OD
END

```

11

What are the different types of EAs

- Historically different flavours of EAs have been associated with different representations
 - Binary strings : **Genetic Algorithms**
 - Real-valued vectors : **Evolution Strategies**
 - Finite state Machines: **Evolutionary Programming**
 - LISP trees: **Genetic Programming**
- These differences are largely irrelevant, best strategy
 - choose representation to suit problem
 - choose variation operators to suit representation
- Selection operators only use fitness and so are independent of representation

12

Examples of EAs – Genetic Algorithms (GAs)

- **Advantages:**


- Easy to code
- Can provide multiple solutions
- Simple ways to avoid local minima/maxima (not guarantee)
- Can be parallelized

- **Disadvantages:**

- They can be slow
- Can be hard to design a good fitness function
- Can be hard to represent solutions of the problem as GA chromosomes

13

Example: the 8 queens problem




Place 8 queens on an 8x8 chessboard in such a way that they cannot attack each other

14

The 8 queens problem: Representation

a board configuration

a permutation of
the numbers 1 - 8



Obvious mapping

1	3	5	2	6	4	7	8
---	---	---	---	---	---	---	---


15

The 8 queens problem: Representation

**Suitable representation?
GENETIC ALGORITHM**

a board configuration

a permutation of
the numbers 1 - 8



Obvious mapping

1	3	5	2	6	4	7	8
---	---	---	---	---	---	---	---

16

Evaluation (Fitness) Function

- Represents the requirements that the population should adapt to
- a.k.a. *quality* function or *objective function*
- Assigns a single real-valued fitness to each phenotype which forms the basis for selection
 - So the more discrimination (different values) the better
- Typically we talk about fitness being maximised
 - Some problems may be best posed as minimisation problems, but conversion is trivial

17

8 Queens Problem: Fitness evaluation

- Penalty of one queen:
the number of queens she can attack.
- Penalty of a configuration:
the sum of the penalties of all queens.
- Note: penalty is to be minimized
- Fitness of a configuration:
inverse penalty to be maximized

18

Population

- Holds (representations of) possible solutions
- Selection operators usually take whole population into account i.e., reproductive probabilities are *relative* to *current* generation
- **Diversity** of a population refers to the number of different **fitnesses** and/or **individuals/chromosomes** present (note: not the same thing)



19

Parent Selection Mechanism

- Assigns variable probabilities of individuals acting as parents depending on their fitnesses
- Usually **probabilistic**
 - high quality solutions more likely to become parents than low quality
 - but not guaranteed
 - even worst in current population usually has non-zero probability of becoming a parent
- This *stochastic* nature can aid escape from local optima



20

Variation Operators

- Role is to generate new candidate solutions
- Usually divided into two types according to their **arity** (number of inputs):
 - Arity = 1 (aka. unary operators): **mutation**
 - Arity > 1 : Recombination operators
 - Arity = 2 (aka. binary operators): typically called **crossover**
- There has been much debate about relative importance of recombination and mutation
 - Nowadays most EAs use both
 - Choice of particular variation operators is representation dependant

21

Recombination


- Merges information from parents into offspring
- Choice of what information to merge is **stochastic**
- Most offspring may be worse, or the same as the parents
- Hope is that some are better by combining elements of genotypes that lead to good traits
- Principle has been used for millennia by breeders of plants and livestock

22

The 8 queens problem: Recombination

Combining two permutations into two new permutations:

- choose random crossover point
- copy first parts into children
- create second part by inserting values from other parent:
 - in the order they appear there
 - beginning after crossover point
 - skipping values already in child



23

Mutation

- Acts on one genotype and delivers another
- Element of **randomness** is essential and differentiates it from other unary heuristic operators
- Importance ascribed depends on representation and dialect:
 - Binary GAs – background operator responsible for preserving and introducing diversity
 - EP for FSM's/ continuous variables – only search operator
 - GP – hardly used
- May guarantee connectedness of search space and hence convergence proofs

24

The 8 queens problem: Mutation

Small variation in one permutation, e.g.:

- swapping values of two randomly chosen positions,



25

Survivor Selection

- a.k.a. **replacement**
- Most EAs use fixed population size so need a way of going from (parents + offspring) to next generation
- Often deterministic
 - Fitness based : e.g., rank parents+offspring and take best
 - Age based: make as many offspring as (reproduced) parents and delete all those parents
- Sometimes do combination (elitism)

26

The 8 queens problem: Selection

- Parent selection:
 - Pick randomly 5 parents and take best two to undergo crossover
- Survivor selection (replacement)
 - When inserting a new child into the population, choose an existing member to replace by:
 - sorting the whole population by decreasing fitness
 - enumerating this list from high to low
 - replacing the first with a fitness lower than the given child

27

Initialisation / Termination

- Initialisation usually done at random,
 - Need to ensure even spread and mixture of possible values
 - Can include existing solutions, or use problem-specific heuristics, to “seed” the population
- Termination condition checked every generation
 - Reaching some (known/hoped for) fitness
 - Reaching some maximum allowed number of generations
 - Reaching some minimum level of diversity
 - Reaching some specified number of generations without fitness improvement

28

8 Queens Problem: Summary


Representation	Permutations
Recombination	"Cut-and-crossfill" crossover
Recombination probability	100%
Mutation	Swap
Mutation probability	80%
Parent selection	Best 2 out of random 5
Survival selection	Replace worst
Population size	100
Number of Offspring	2
Initialisation	Random
Termination condition	Solution or 10,000 fitness evaluation

Note that this is ***only one possible*** set of choices of operators and parameters


29

Typical behaviour of an EA


Phases in optimising on a 1-dimensional fitness landscape



Early phase:
quasi-random population distribution




Mid-phase:
population arranged around/on hills



Late phase:
population concentrated on high hills

30

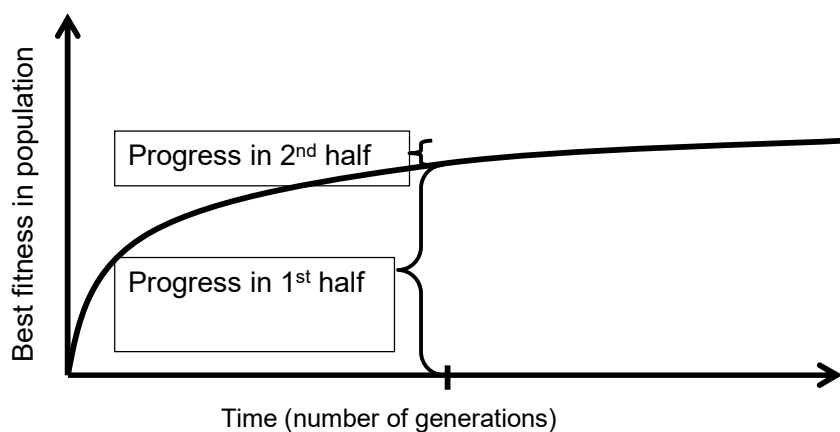
Typical run: progression of fitness



Typical run of an EA shows so-called “anytime behavior”

31


Are long runs beneficial?



- **Answer:**
 - it depends how much you want the last bit of progress
 - it may be better to do more shorter runs

32

Is it worth expending effort on smart initialisation?



- **Answer** : it depends:
 - possibly, if good solutions/methods exist.
 - care is needed


33

Evolutionary Algorithms in Context

- There are many views on the use of EAs as robust problem solving tools
- For most problems a problem-specific tool may:
 - perform better than a generic search algorithm on most instances,
 - have limited utility,
 - not do well on all instances
- Goal is to provide robust tools that provide:
 - evenly good performance
 - over a range of problems and instances

34

EAs as problem solvers: Goldberg's 1989 view




35

EAs and domain knowledge

- Trend in the 90's:
adding problem specific knowledge to EAs
(special variation operators, repair, etc)
- Result: EA performance curve "deformation":
 - better on problems of the given type
 - worse on problems different from given type
 - amount of added knowledge is variable
- Recent theory suggests the search for an "all-purpose" algorithm may be fruitless

36

Michalewicz' 1996 view



37

EC and Global Optimisation

- Global Optimisation: search for finding best solution x^* out of some fixed set S
- Deterministic approaches
 - e.g. box decomposition (branch and bound etc)
 - Guarantee to find x^* , but may run in super-polynomial time
- Heuristic Approaches (generate and test)
 - rules for deciding which $x \in S$ to generate next
 - no guarantees that best solutions found are globally optimal

38

Applicable situations

- Often used for optimization (scheduling, design, etc.) problems, though can be used for many other things as well, as we'll see a bit later.
 - Good problem for EAs: Scheduling air traffic
 - Bad problems for EA: Finding large primes (why?), 2D pathfinding (why?)

39

39

Applicable situations

- EAs work best when the “fitness landscape” is continuous (in some dimensions). This is also true of standard search, e.g. A*.
 - Intuitively, this just means that we can find a heuristic that gives a rough idea of how close a candidate is to being a solution.





Image source: scholarpedia.org

40

40

Examples - EA in the wild

- Image compression – evolving the Mona Lisa



<http://rogersaling.com/2008/12/07/genetic-programming-evolution-of-mona-lisa/>

41

41

Evolving the Mona Lisa

- Uses only 50 polygons of 6 vertices each.
- Population size of 1, no crossover – parent compared with child, and superior image kept.
- Assuming each polygon has 4 bytes for color (RGBA) and 2 bytes for each of 6 vertices, this image only requires 800 bytes.
- However, compression time is prohibitive and storage is cheaper than processing time. ☹️

42

42

More evolved images



<http://rogersaling.com/2008/12/07/genetic-programming-evolution-of-mona-lisa/>

43